

UNIVERSITY COLLEGE LONDON  
DEPARTMENT OF SPACE AND CLIMATE PHYSICS

**Candidate Code:** HYXC3

Programme Title: MSc Scientific Computing

Module Code: SPCE0038

**Module Title: Machine Learning with Big Data**

## End Assessment

In submitting this coursework, I assert that the work presented is entirely my own except where properly marked and cited.

Date of Submission:	11/05/20
------------------------	----------

## Question 1

### 1(a)

With reference to the diagram of the basic *logistic unit* on the following page:

The input vector  $\mathbf{x}$  is the input to the *logistic unit*.

Each input  $x_i$  has an associated weight  $\theta_i$ . The weights are set to a random value prior to *training*. The *training* process determines these weights.

The product of each input  $x_i$  and weight  $\theta_i$  is summed to produce a weighted sum  $z$ .

The output,  $h_\theta(x)$ , is the the non-linear *activation function*,  $h$ , applied to  $z$ .

### 1(b)

Consider the diagram of Question 1(a).

The weighted sum of the inputs,  $z$ , is:

$$z = \sum_{j=1}^n \theta_j x_j = \theta^T x \quad (1)$$

where  $x_i$  is the  $i^{th}$  element of input vector  $\mathbf{x}$  of length  $n$ , and  $\theta_i$  is the associated *weight*.

And, the output from the *logistic unit*,  $h_\theta(x)$ , is:

$$h_\theta(x) = h(z) \quad (2)$$

where  $h$  is a non-linear *activation function*.

### Question 1(a) – Basic Logistic Unit



$x_i$ : Input

$\theta_i$ : Weight

Weighted Sum:  $z = \sum_{j=1}^n \theta_j x_j = \theta^T x$

Activation:  $a = h(z)$

### 1(c)

See diagram on following page.

### 1(d)

Consider the diagram of Question 1(c).

Firstly, consider the data transformation from the input vector  $\mathbf{x}$  to the hidden layer. We now need two indices, one for the input vector elements, and one for the *hidden layer* nodes. We will use  $i$  for the input vector index, and  $j$  for the *hidden layer* nodes.

The weighted sum of the inputs at the  $j^{\text{th}}$  *hidden layer* node is:

$$z_j = \sum_{i=1}^n \theta_{ij} x_i \quad (3)$$

where  $\theta_{ij}$  is the weight between input element  $i$  and *hidden layer* node  $j$ , and  $n$  is the length of the input vector  $\mathbf{x}$ .

Secondly, the output from each *hidden layer* node is the non-linear *activation function*,  $h$ , applied to each  $z_j$ :

$$h_{\theta_j}(x) = h(z_j) \quad (4)$$

And finally, the output from the whole network,  $h_{\Theta}(x)$  is the sum of the *hidden layer* outputs:

$$h_{\Theta}(x) = \sum_{j=1}^m h_{\theta_j}(x) \quad (5)$$

where  $m$  is the number of *hidden layer* nodes.

# Question 1(c) – Fully Connected, Feed Forward, Artificial Neural Network



$x_i$ : Input

$\theta_{ij}$ : Weight, e.g.  $\theta_{11}$   $\theta_{33}$

Weighted Sums: 
$$z_j = \sum_{i=1}^n \theta_{ij} x_i$$

Activations: 
$$a_j = h(z_j)$$

- Input Layer Logistic Units
- Hidden Layer Logistic Units
- Output Node

## 1(e)

The cost function typically used to train neural networks for regression problems is *mean square error*:

$$MSE(\Theta) = \frac{1}{m} \sum_i \sum_j (p_j^{(i)} - y_j^{(i)})^2 \quad (6)$$

The cost function typically used to train neural networks for classification problems is *cross-entropy*:

$$C(\Theta) = -\frac{1}{m} \sum_i \sum_j y_j^{(i)} \log(p_j^{(i)}) \quad (7)$$

## 1(f)

Artificial Neural Networks (ANNs) are described as *shallow* or *deep*, and *wide* or *narrow*. *shallow* or *deep* refers to the number of layers in the network, and *wide* or *narrow* refers to the number of nodes in each layer.

The *credit assignment path*, the CAP, of a neural network is a measure of the number of data transformations that occur as data passes through the network. For *feed-forward* networks the CAP is the number of *hidden layers* plus one.

A *deep* neural network is generally considered to be a network with multiple layers and a CAP > 2.

## 1(g)

The *universal approximation theorem* states that, with appropriate parameters, single hidden layer feed-forward neural networks are *universal approximators*. This means they can represent any continuous function. However, this requires an exponentially larger number of hidden layer nodes. And, training will not necessarily determine the parameters.

Deep networks provide a powerful representational framework because they have the potential to be *universal approximators*, but with a limited width of hidden nodes. This makes the implementation of *universal approximators* more feasible.

## Question 2

### 2(a)

*Gradient Descent* algorithms attempt to find the parameters  $\theta$  which minimise the cost function  $C(\theta)$  over the *training set* using an iterative process:

$$\theta^{i+1} = \theta^i - \alpha \nabla_{\theta} C(\theta) \quad (8)$$

where  $\alpha$  is the *learning rate*.

*Batch Gradient Descent* uses the entire *training set* at each iteration to calculate the gradient partial derivatives,  $\nabla_{\theta} C(\theta)$ . This produces accurate values for the partial derivatives, but is potentially slow for large *training sets*.

*Stochastic Gradient Descent* uses a random sub-set of the *training set* at each iteration to calculate the gradient partial derivatives,  $\nabla_{\theta} C(\theta)$ . This is faster than *Batch Gradient Descent*, but can produce erratic values for the partial derivatives.

For *convex* cost functions *Batch Gradient Descent* will always converge to the *local minima* which is also the *global minima*. This is not the case for *non-convex* cost functions, where *local minima* are not necessarily *global minima*. A bad choice of  $\theta^0$  may result in *Batch Gradient Descent* getting “stuck” in a *local minima*. Because the gradient partial derivatives of *Stochastic Gradient Descent* are erratic, this provides a mechanism of “jumping out of” a *local minima* and improves the probability of finding the *global minima*.

### 2(b)

When attempting to find the minimum of a cost function using *Stochastic Gradient Descent*, the iterative process “jumps” around the minimum, and it is difficult to determine when a minimum has been reached. For this reason alternative optimisation algorithms are typically considered for training.



## 2(c)

Consider the iterative minimisation process:

$$\theta^{i+1} = \theta^i - \alpha \nabla_{\theta} C(\theta) \quad (9)$$

where  $\alpha$  is the *learning rate*.

At each step the process “advances” towards the minimum by the step size  $-\alpha \nabla_{\theta} C(\theta)$ , which uses the *current* gradient.

The *Momentum Optimisation Algorithm* introduces the idea of including *previous* gradients into the step size. At each step the *current* gradient is summed into a momentum term,  $m$  (remember we use the negative gradient):

$$m^{i+1} = \beta m^i - \alpha \nabla_{\theta} C(\theta) \quad (10)$$

The  $\beta$  parameter is used to prevent the momentum getting too large, and is set between 0 and 1, typically 0.9.

The momentum  $m$  is then used to update  $\theta$ :

$$\theta^{i+1} = \theta^i + m \quad (11)$$

The result is that we now have an *acceleration* towards the minimum.

**2(d)**

**2(e)**

**2(f)**

**2(g)**

**2(h)**

### Question 3

3(a)

TODO

3(b)

TODO

3(c)

TODO

3(d)

TODO

3(e)

TODO

3(f)

TODO

## Question 4

4(a)

TODO

4(b)

TODO

4(c)

TODO

4(d)

TODO

4(e)

TODO

4(f)

TODO

## question\_4f

May 7, 2020

```
[ ]: # Fetch batch function:

def fetch_batch(epoch, batch_index, batch_size):

    return X_batch, y_batch

# Set up computational graph:

import tensorflow as tf
reset_graph ()

n_epochs = 1000
learning_rate = 0.01

X = tf.constant(scaled_housing_data_plus_bias, dtype=tf.float32, name="X")
y = tf.constant(housing_data_target, dtype=tf.float32, name="y")

theta = tf.Variable(tf.random_uniform([n + 1, 1], -1.0, 1.0), name="theta")
y_pred = tf.matmul(X, theta, name="predictions")
error = y_pred - y
mse = tf.reduce_mean(tf.square(error), name="mse")
optimizer = tf.train.GradientDescentOptimizer(learning_rate)
training_op = optimizer.minimize(mse)

# Execute:

init = tf.global_variables_initializer()

with
tf.Session() as sess:
    sess.run(init)
    for epoch in range(n_epochs):
        if epoch % 100 == 0:
            print("Epoch", epoch, "MSE=", mse.eval()) sess.run(training_op)
    best_theta = theta.eval()
```

```

1  # Fetch batch function:
2
3  def fetch_batch(epoch, batch_index, batch_size):
4      return X_batch, y_batch
5
6
7  # Set up computational graph:
8
9  import tensorflow as tf
10 reset_graph ()
11
12 n_epochs = 1000
13 learning_rate = 0.01
14
15 X = tf.constant(scaled_housing_data_plus_bias, dtype=tf.float32, name="X")
16 y = tf.constant(housing_data_target, dtype=tf.float32, name="y")
17
18 theta = tf.Variable(tf.random_uniform([n + 1, 1], -1.0, 1.0), name="theta")
19 y_pred = tf.matmul(X, theta, name="predictions")
20 error = y_pred - y
21 mse = tf.reduce_mean(tf.square(error), name="mse")
22 optimizer = tf.train.GradientDescentOptimizer(learning_rate)
23 training_op = optimizer.minimize(mse)
24
25
26 # Execute:
27
28 init = tf.global_variables_initializer()
29
30 with tf.Session() as sess:
31     sess.run(init)
32     for epoch in range(n_epochs):
33         if epoch % 100 == 0:
34             print("Epoch", epoch, "MSE=", mse.eval())
35             sess.run(training_op)
36     best_theta = theta.eval()

```

Listing 1: Question 4f

## Question 5

5(a)

TODO

5(b)

TODO

5(c)

TODO

5(d)

TODO

5(e)

TODO

5(f)

TODO